

# AI504: Programming for Artificial Intelligence

## Week 1: Introduction

Edward Choi

Grad School of AI

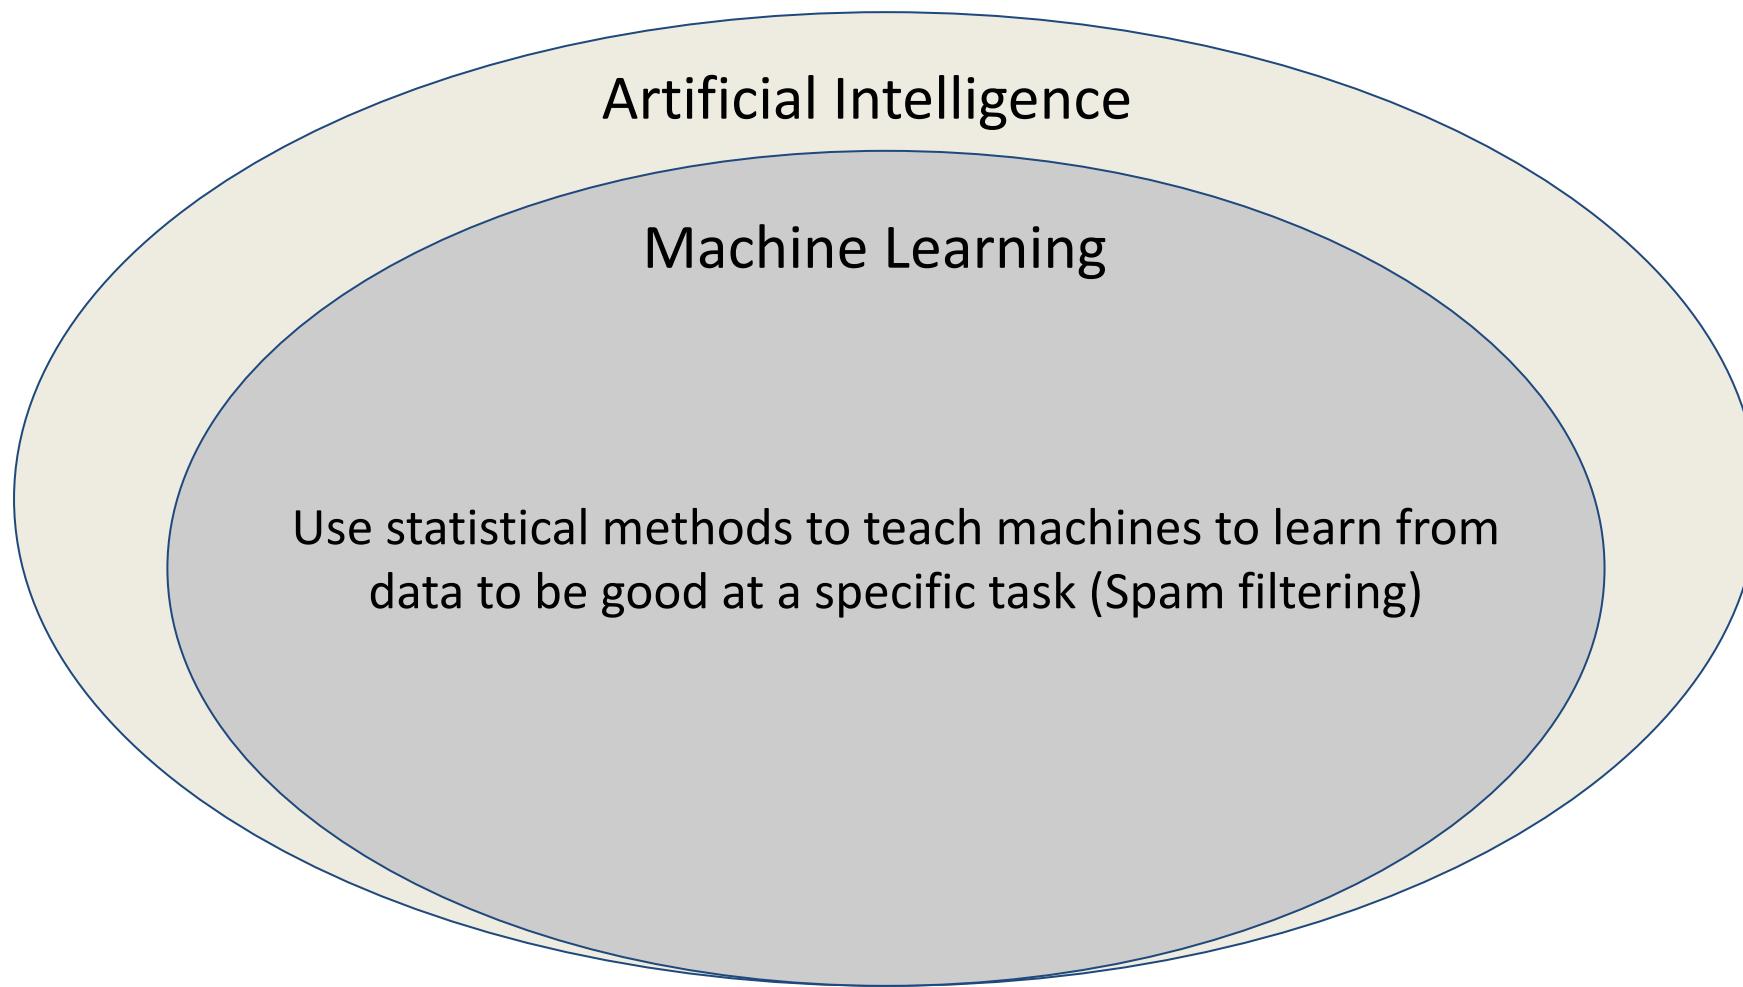
[edwardchoi@kaist.ac.kr](mailto:edwardchoi@kaist.ac.kr)

# What is AI?

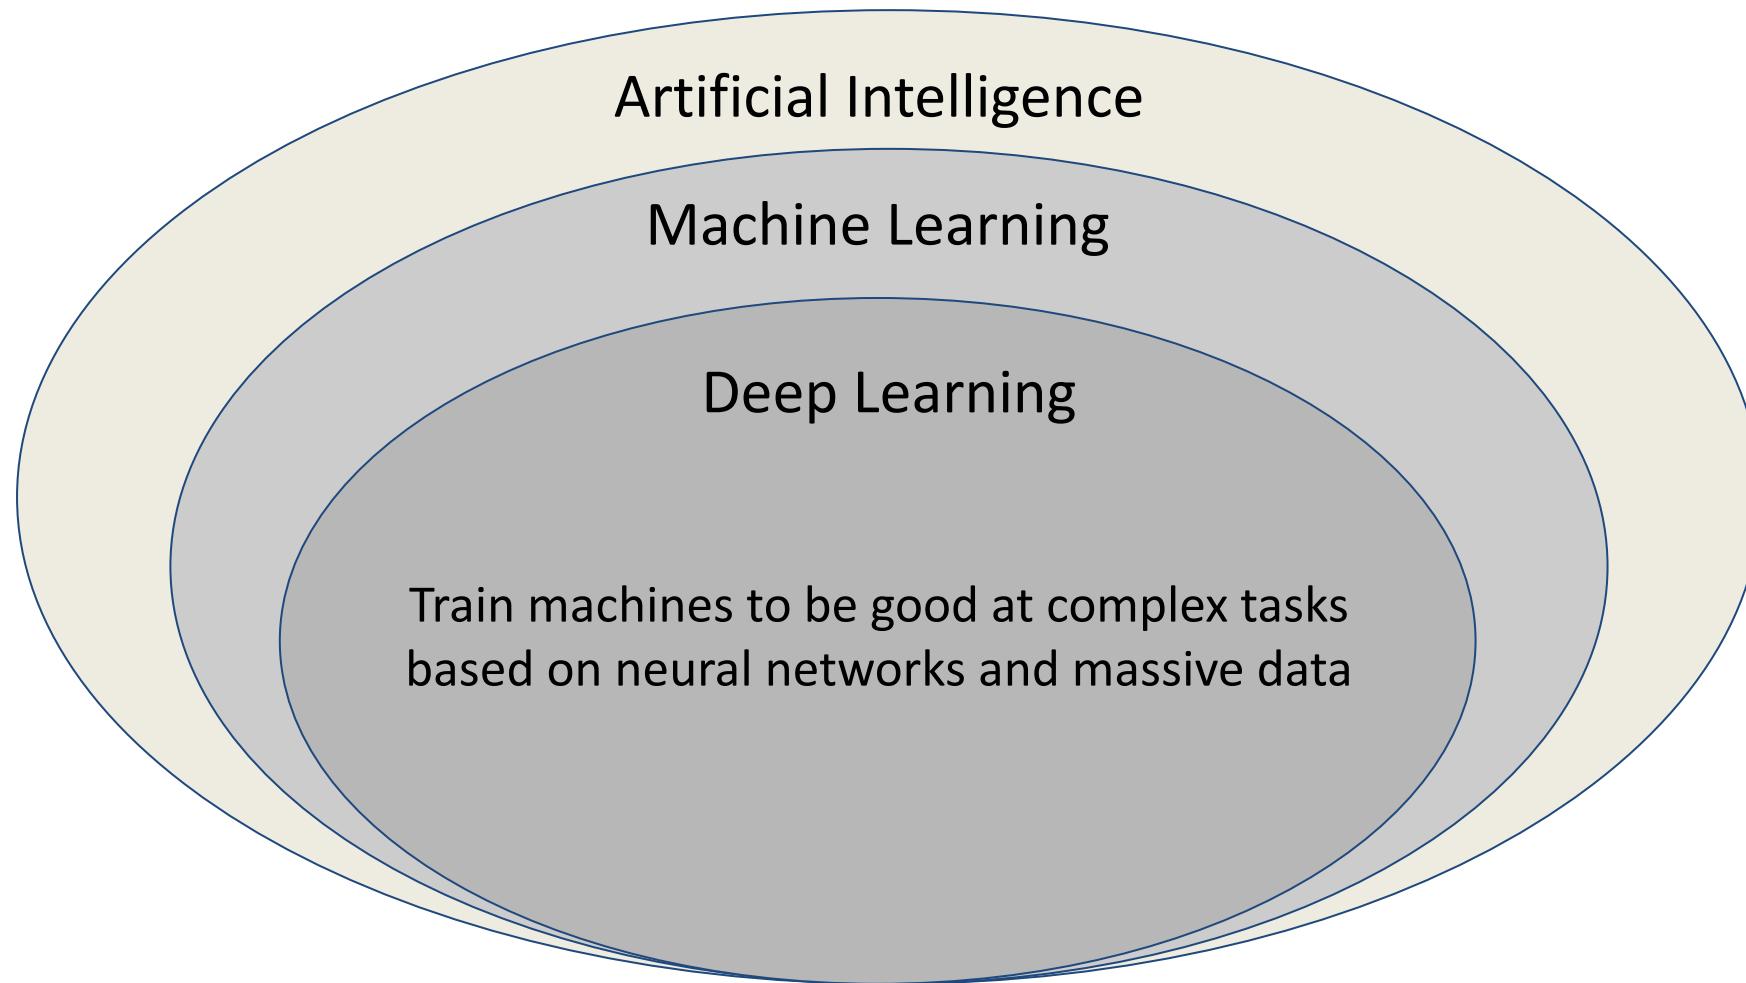
Artificial Intelligence

Make machines/computers mimic human intelligence  
Concept as old as the computer (Chess program by Alan Turing)

# What is Machine Learning?

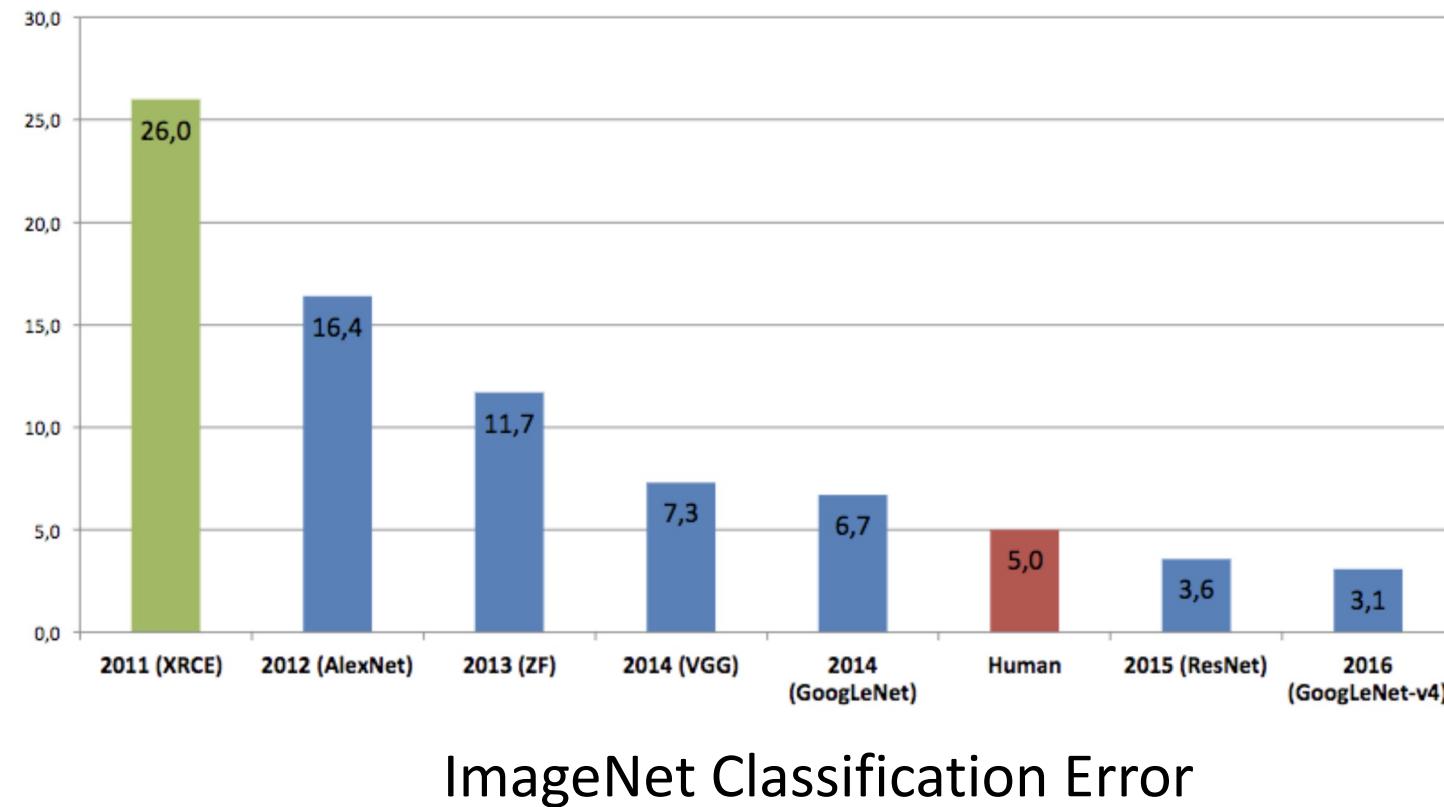


# What is Deep Learning?



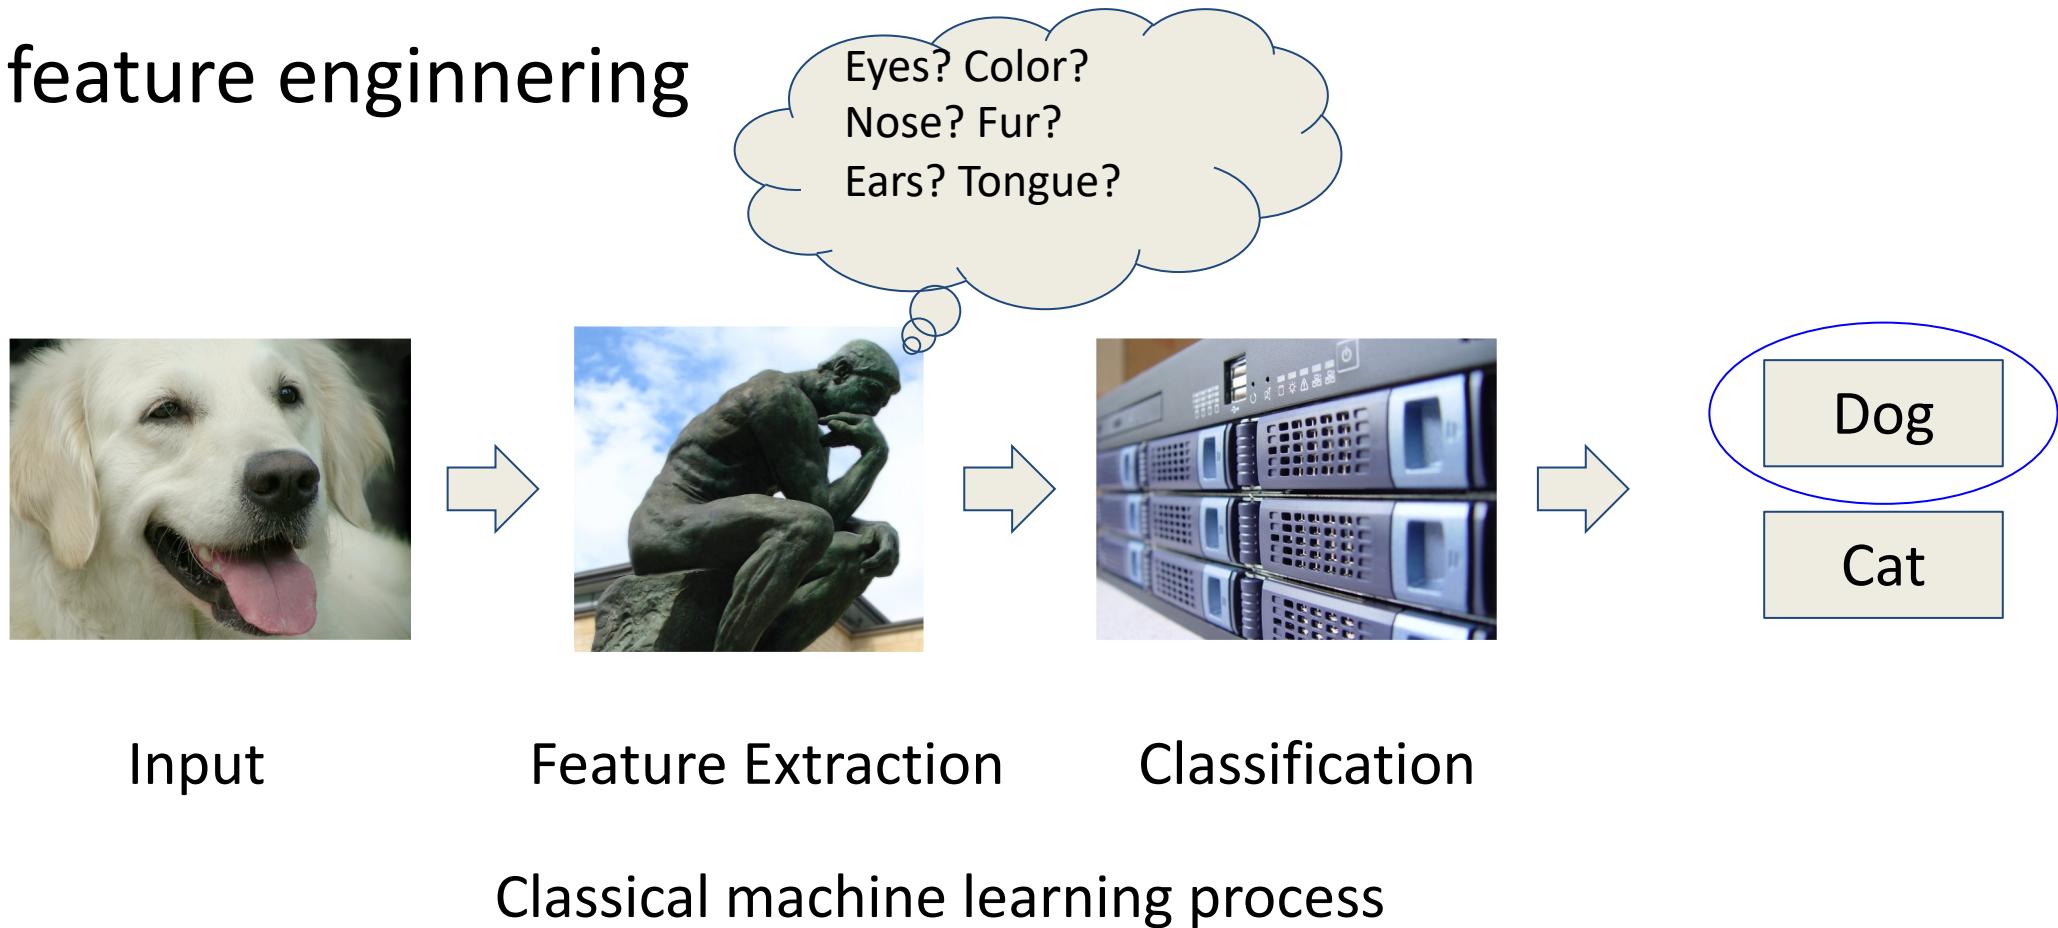
# Why Deep Learning?

State-of-the-art performance



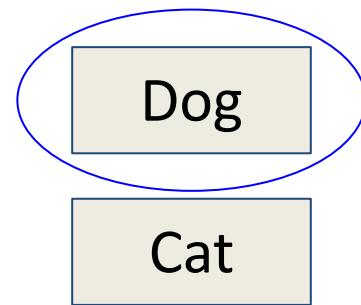
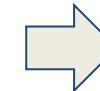
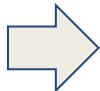
# Why Deep Learning?

Less feature engineering



# Why Deep Learning?

Less feature engineering



Input

Feature Extraction + Classification

Deep learning process

# How is this Possible?

Large data

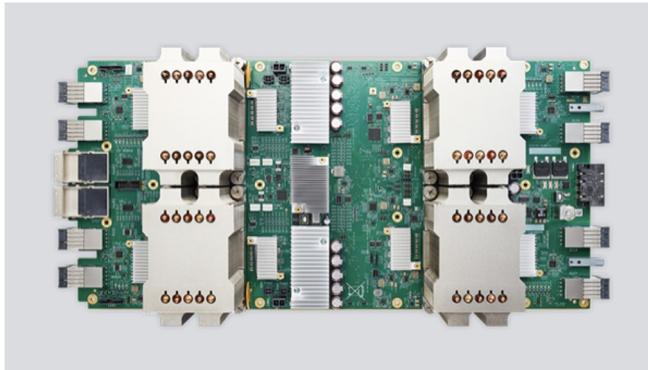
- Social network services
  - Youtube, Instagram, Twitter
- Collective intelligence
  - Wikipedia
- Mass media
  - News articles



ImageNet

# How is this Possible?

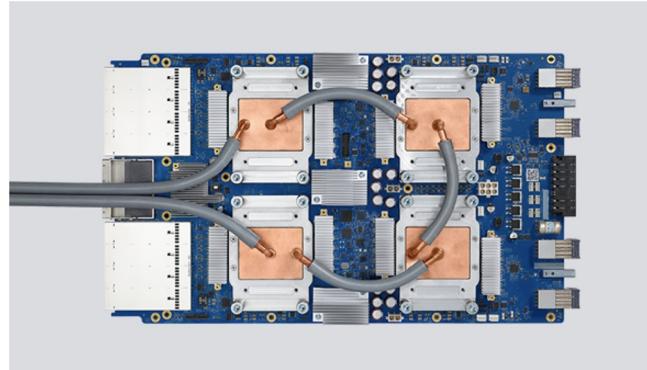
Large data + Powerful machines



**Cloud TPU v2**

180 teraflops

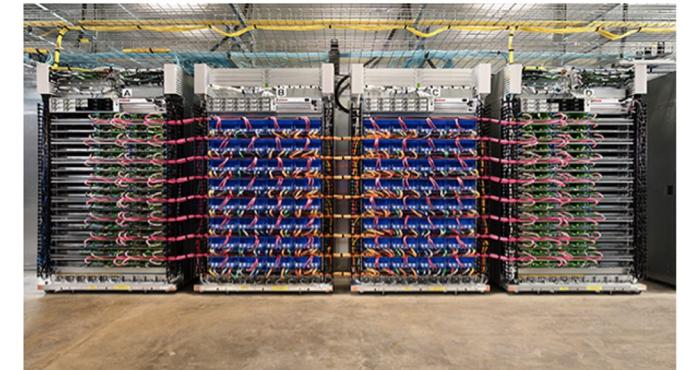
64 GB High Bandwidth Memory (HBM)



**Cloud TPU v3**

420 teraflops

128 GB HBM



**Cloud TPU v2 Pod (alpha)**

11.5 petaflops

4 TB HBM

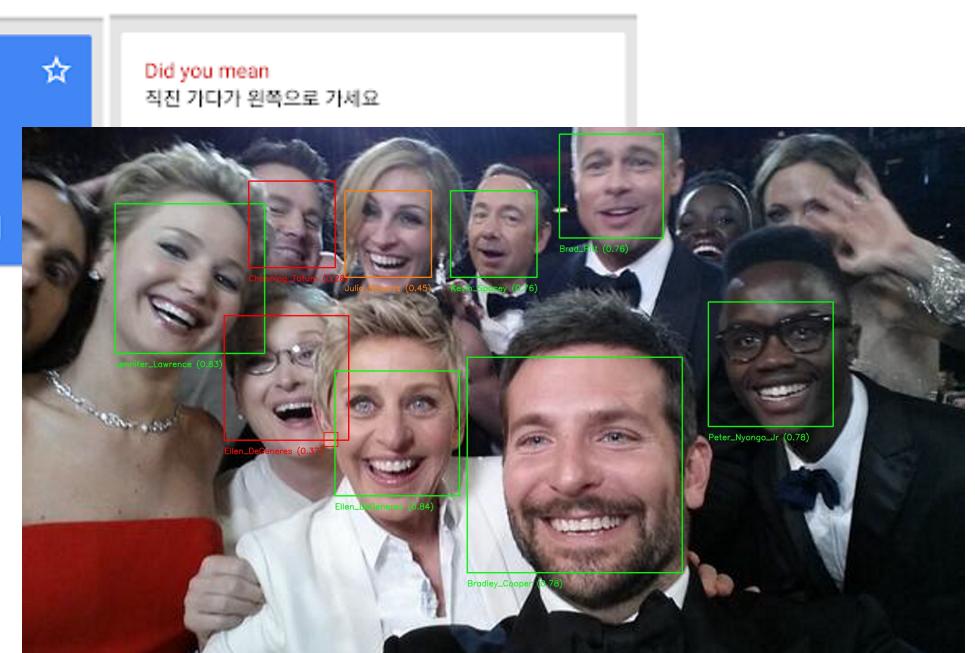
2-D toroidal mesh network

# Modern AI



Two screenshots of the Google Translate mobile application. Both screens have a blue header with the Google Translate logo and a menu icon. The left screen shows "English" as the source language and "Korean" as the target language. The right screen shows "Korean" as the source language and "English" as the target language. Both screens show the text "where is the train station" and its Korean translation "직진 하다가 왼쪽으로 가세요".

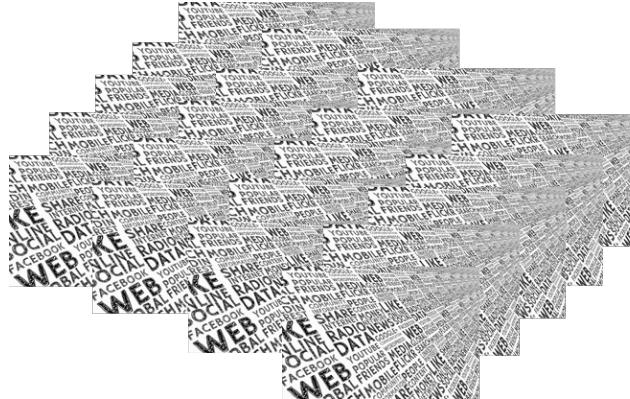
Two screenshots of the Google Translate mobile application. Both screens have a blue header with the Google Translate logo and a menu icon. The left screen shows "Korean" as the source language and "English" as the target language. The right screen shows "English" as the source language and "Korean" as the target language. Both screens show the text "Did you mean" followed by the Korean phrase "직진 기다기 왼쪽으로 가세요" and its English translation "jigjin hadaga oenjjog-eulo gaseyo".



# (Almost) Infinite Compute + (Almost) Infinite Data



Probably >1000 V100's



570GB Text (400 Billion Tokens)



175B Parameters  
\$4~5 Million for Training

# GPT3

- Open AI's unsupervised generative language model
  - Using 96 layers of Transformer
  - BERT uses 24 layers

Title: United Methodists Agree to Historic Split

Subtitle: Those who oppose gay marriage will form their own denomination

Article: After two days of intense debate, the United Methodist Church has agreed to a historic split - one that is expected to end in the creation of a new denomination, one that will be "theologically and socially conservative," according to The Washington Post. The majority of delegates attending the church's annual General Conference in May voted to strengthen a ban on the ordination of LGBTQ clergy and to write new rules that will "discipline" clergy who officiate at same-sex weddings. But those who opposed these measures have a new plan: They say they will form a separate denomination by 2020, calling their church the Christian Methodist denomination.

The Post notes that the denomination, which claims 12.5 million members, was in the early 20th century the "largest Protestant denomination in the U.S.," but that it has been shrinking in recent decades. The new split will be the second in the church's history. The first occurred in 1968, when roughly 10 percent of the denomination left to form the Evangelical United Brethren Church. The Post notes that the proposed split "comes at a critical time for the church, which has been losing members for years," which has been "pushed toward the brink of a schism over the role of LGBTQ people in the church." Gay marriage is not the only issue that has divided the church. In 2016, the denomination was split over ordination of transgender clergy, with the North Pacific regional conference voting to ban them from serving as clergy, and the South Pacific regional conference voting to allow them.

A "whatpu" is a small, furry animal native to Tanzania. An example of a sentence that uses the word whatpu is:

We were traveling in Africa and we saw these very cute whatpus.

To do a "farduddle" means to jump up and down really fast. An example of a sentence that uses the word farduddle is:

One day when I was playing tag with my little sister, she got really excited and she started doing these crazy farduckles.

A "yalubalu" is a type of vegetable that looks like a big pumpkin. An example of a sentence that uses the word yalubalu is:

I was on a trip to Africa and I tried this yalubalu vegetable that was grown in a garden there. It was delicious.

A "Burringo" is a car with very fast acceleration. An example of a sentence that uses the word Burringo is:

In our garage we have a Burringo that my father drives to work every day.

A "Gigamuru" is a type of Japanese musical instrument. An example of a sentence that uses the word Gigamuru is:

I have a Gigamuru that my uncle gave me as a gift. I love to play it at home.

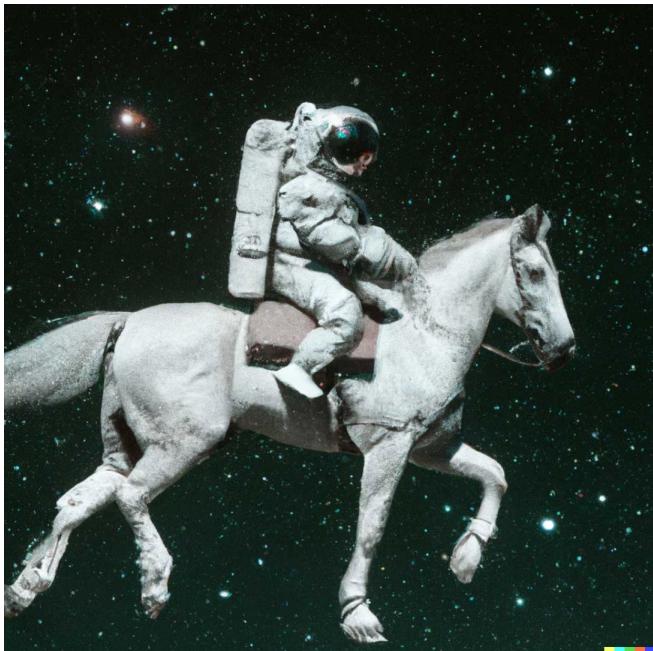
To "screeg" something is to swing a sword at it. An example of a sentence that uses the word screeg is:

We screeghed at each other for several minutes and then we went outside and ate ice cream.

# DALL-E 2, Imagen, Parti

- Open AI & Google's conditional text-to-image generator
  - Transformer or Diffusion

An astronaut riding a horse in a photorealistic style



A chrome-plated duck with a golden beak arguing with a golden turtle in a forest

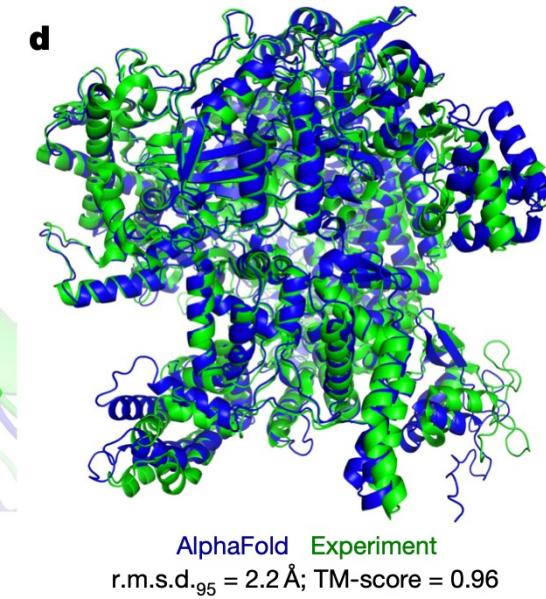
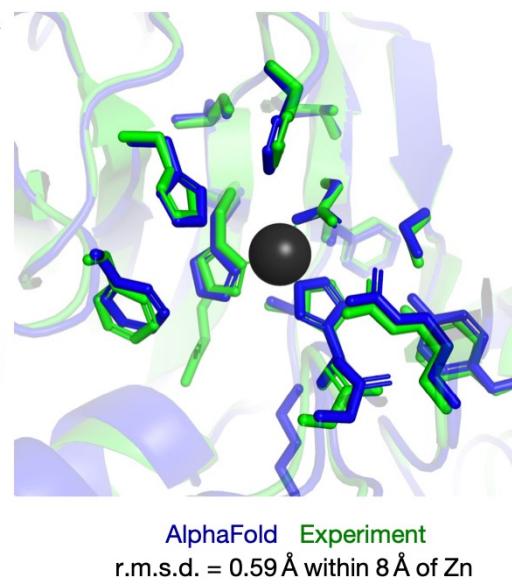
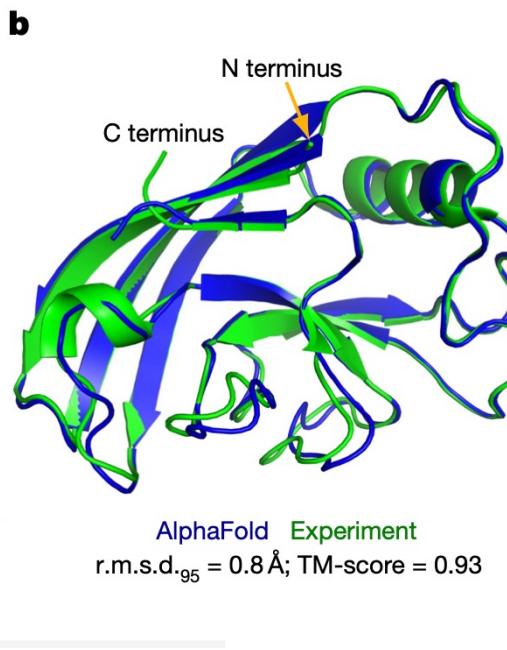
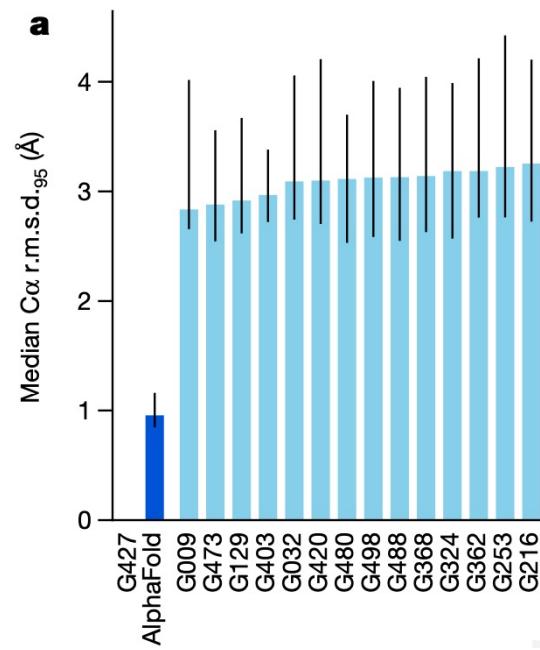


A photo of an Athenian vase with a painting of pangolins playing basketball in the style of Egyptian hieroglyphics.

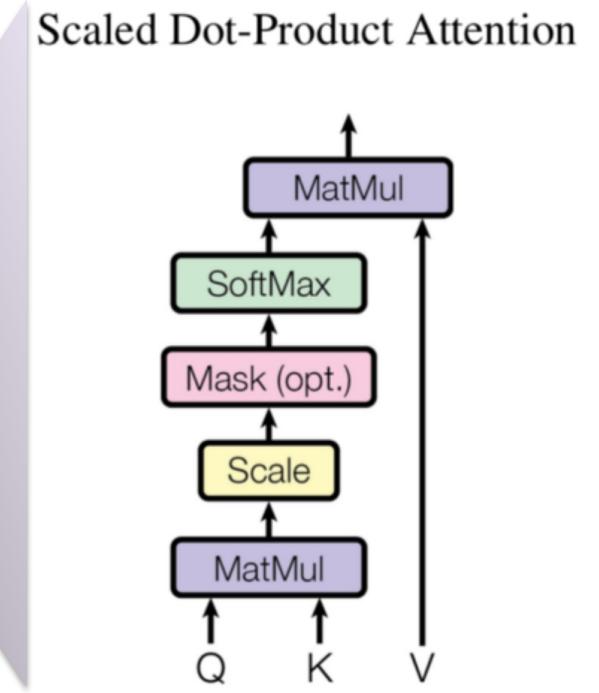
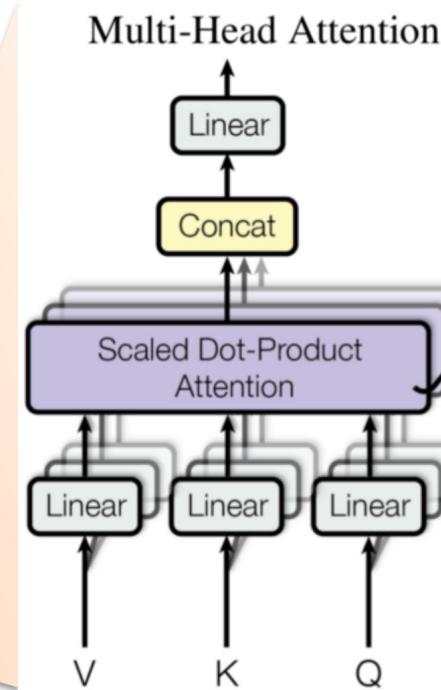
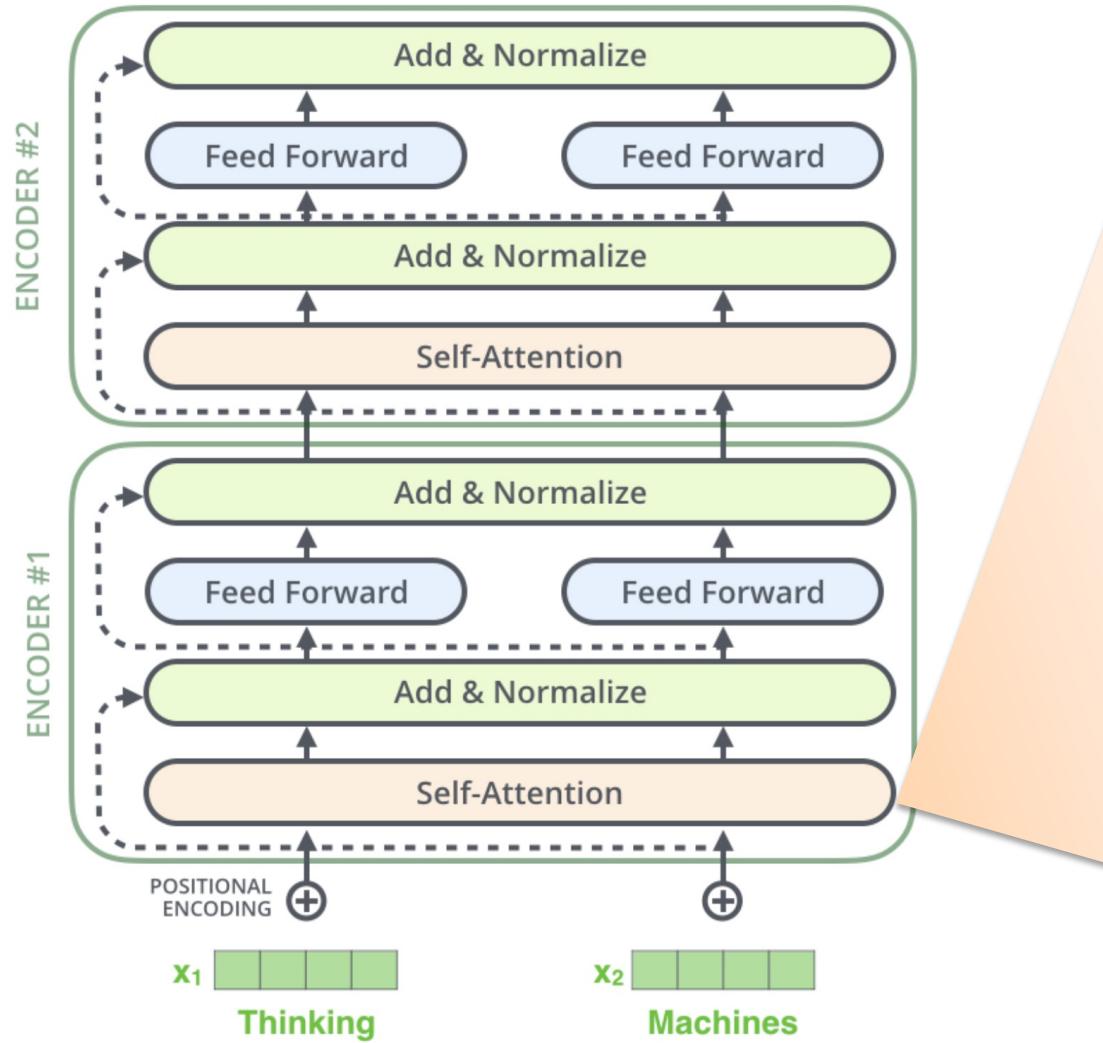


# AlphaFold2

- DeepMind's protein structure prediction model
  - Using 48 layers of Evoformer
    - Special self-attention for graphs in 3D



# Transformer Architecture



# Goal

- Learn to build deep learning models.
  - So that you can replicate papers
  - So that you can realize your ideas
  - So that you can conduct AI research
- This course teaches only the very basics.
  - Practice makes perfect!

# Structure

- Schedule:
  - Lecture on Tuesday
    - Conducted by the lecturer
  - Practice on Thursday
    - Conducted by the TA
- Assignments
  - Project 1
    - Image synthesis using GAN
  - Project 2
    - French-English Translation using Transformers
- Grading
  - Letter grade
  - Project 1: 45%, Project 2: 45%, Attendance: 10%

# Weekly Plan

1. Intro + Numpy
2. Basic Machine Learning + Scikit-learn
3. PyTorch Intro + Logistic Regression + Multi-layer Perceptron
4. Autoencoders (& Denoising Autoencoders)
5. Variational Autoencoders
6. Generative Adversarial Networks
7. Convolutional Neural Networks
8. Midterm Week (Project 1)
9. Word2Vec + Subword Encoding
10. Recurrent Neural Networks & Sequence-to-Sequence
11. Transformers
12. BERT (& GPT)
13. Image-Text Multimodal Learning
14. Deep Diffusion Probabilistic Model
15. Graph Neural Networks
16. Final Week (Project 2)

# Material & Discussion

- No textbook
- Materials will be posted on Classum if necessary.
- Classum
  - Join as Participant Link: [www.classum.com/ej4fmx](http://www.classum.com/ej4fmx)

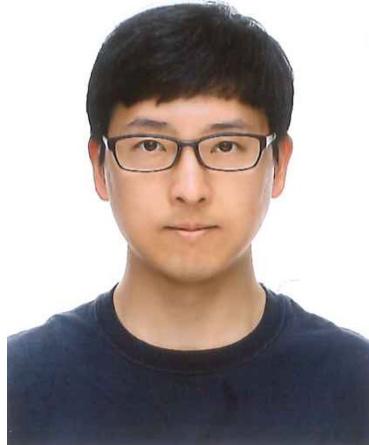
# Expectation

- Lecture: You are expected to attend
- Project
  - Individual effort
  - Deliverables:
    - Link to the model weights
    - Your model will be evaluated automatically
      - Project 1: FID score
      - Project 2: BLEU score

# Teaching Team

- Lecturer
  - Edward Choi
    - [edwardchoi@kaist.ac.kr](mailto:edwardchoi@kaist.ac.kr)
    - <https://mp2893.com>
- TA
  - TBA

# Edward Choi, 최윤재

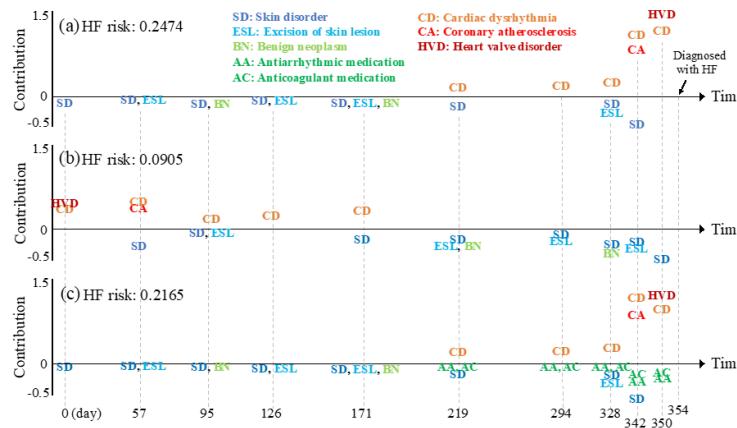


- Education
  - Ph.D. (computer science), Georgia Tech, 2014-2018
    - Thesis: Interpretable deep learning for longitudinal electronic health records
- Professional Experience
  - ETRI (2010-2014)
  - Sutter Health (2015, 2016)
  - DeepMind & Google (2017)
  - Google Brain & Google Health (2018-2020)
- Research Area
  - Machine Learning, Healthcare, NLP, Multi-modal

# Research Topics

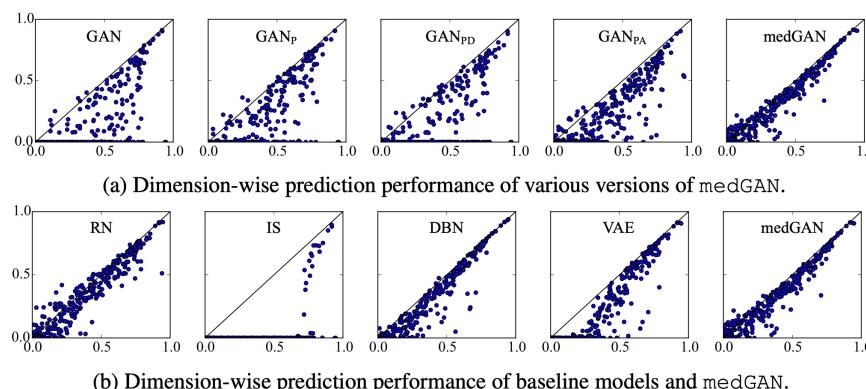
## Interpretable Healthcare Prediction

- Predict medical events (e.g. disease)
- Provide explanation for the prediction
- Recent focus on scaling up



## Synthesizing Healthcare Data

- Healthcare data is hard to share due to privacy risk
- High-quality synthetic samples can accelerate research

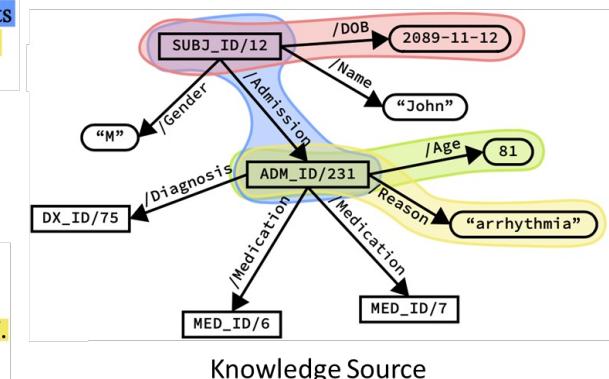


## Question Answering with EHR

- Exploit clinical facts with human language
- Perform complex reasoning with clinical facts
- Have conversation with the AI agent

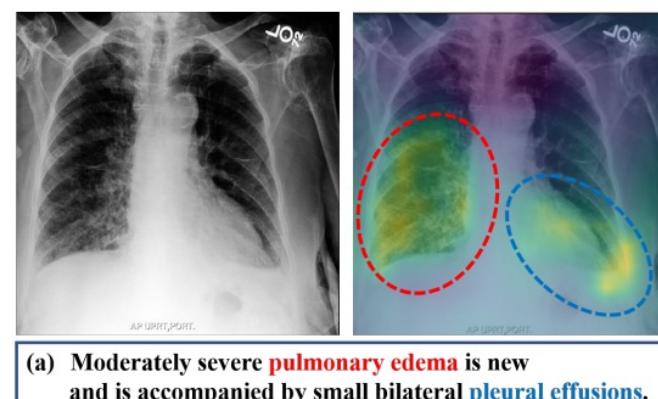
Calculate the **maximum age** of patients who **have arrhythmia primary disease** and were **born after the year 2080**?

A.I.  
SELECT ( MAX ( ?age ) as ?agg )  
WHERE { ?subj\_id </Admission> ?adm\_id .  
?adm\_id </Age> ?age .  
?adm\_id </Reason> "Arrhythmia" .  
?subj\_id </DOB> ?dob .  
filter ( ?dob > 2080-01-01| ) . }



## Multi-modal Learning

- Healthcare data consists of multiple modalities
- Must combine all information to accurately treat patients



# First Assignment

- Install Anaconda
  - Python package for data science
  - Includes Jupyter, Numpy, Scikit-Learn, TensorFlow, PyTorch
  - <https://www.anaconda.com/products/individual>
- All practice sessions will be conducted with Google Colab
  - Python Notebook on the web
  - Can train models using Google's GPU/TPU
  - Session-based (why you need Anaconda)

# Questions?

# AI504: Programming for Artificial Intelligence

## Week 1: Introduction

Edward Choi

Grad School of AI

[edwardchoi@kaist.ac.kr](mailto:edwardchoi@kaist.ac.kr)